

Toward NT0

Single-Backbone Multi-Expert Architecture

Qwen2.5-VL · Point-Cloud Expert · Action Flow-Matching Expert

Team members: Mohammad, Andri, and Artem

Core Idea – VLM Expert

The frozen VLM expert is the **actual** `Qwen2_5_VLDecoderLayer` – not a conversion or approximation.

It retains its native:

- 3D **MRoPE**
- **RMSNorm**
- **GQA** (16 Q / 2 KV)
- Causal + packed-sequence attention mask
- Sliding-window layers

VLM runs at **every** backbone layer. Its K/V (post-RoPE) are extracted from `DynamicCache` and passed to downstream experts with `.detach()`.

PC & Action Experts

`ExpertBlock` = downsized transformer block that matches the VLM’s architectural primitives at reduced width.

Primitive	Status
RMSNorm	✓ shared
GatedMLP	✓ shared
GQA	✓ shared
MRoPE	✓ shared

Initialized from VLM weights by **slicing the first N heads** and **truncating MLP dimensions**.

Placement Modes

Experts are placed at configurable layers (`expert_layer_ids`). Two modes control behavior at *non-expert* layers:

Dense

Experts exist at **every** layer.

- **Cross-attention to VLM**: only at `expert_layer_ids`
- **Self-attention only**: at other layers

Full-depth processing with selective cross-modal flow.

Sparse

Experts exist **only** at `expert_layer_ids`.

- Hidden states **pass through unchanged** at other layers

Trainable params proportional to number of expert layers.

Conditioning Paths

Each signal enters the model through **one or two paths**:

Signal	Token sequence	adaLN (Action only)
Camera images (×K)	VLM: <code><vision_start>...<vision_end></code> per cam	—
Task instruction	VLM: text <code><task>...</task></code>	—
Robot description	VLM: text <code>Robot: ...</code>	—
Embodiment ID	PC + Action: <code>SharedEmbodimentEmbedding</code> → <code>embod_pc</code> , <code>embod_action</code>	—
Attention sinks	PC: <code>sink_pc</code> ; Action: <code>sink_action</code> (separate <code>nn.Parameter</code> s at position 0)	—
Point clouds (×K)	PC: <code><pc_(wrist_)start>...<pc_(wrist_)end></code> per cam	—
Alignment flag	PC: <code><align></code> → <code>Embedding(2)[0 or 1]</code>	—

Conditioning – Proprio

Proprio → Action-stream tokens

- Each proprio reading (current, or current + history) is projected by a single shared `Linear(proprio_dim, d_action)` in `ProprioEncoder`.
- Wrapped in `<proprio_start>` / `<proprio_end>` learned delimiters and prepended to the action chunk inside the Action expert.
- A unified causal mask covers `[<proprio_start>, proprio, <proprio_end>, <action_start>, action, <action_end>]` so action tokens attend back to proprio while staying causal among themselves; the delimiters give the model a learned cue for the proprio block boundary even when P varies sample to sample.

Conditioning – Embodiment

Embodiment → shared bottleneck across PC + Action

- One `nn.Embedding(NUM_ROBOTS=3, d_emb)` base + two zero-init `Linear(d_emb, d_*, bias=False)` heads in `SharedEmbodimentEmbedding`.
- `embodiment_pc` lands at **position 1 of the PC segment** (`sink_pc` occupies position 0); `embodiment_action` at position 1 of the Action segment.
- Both losses train the same base via the two projections; Action loss also reaches `to_pc` indirectly through the no-detach PC→Action cross-attention.
- The frozen VLM still gets per-robot context via the **textual** `Robot: {description}` line – no learned token in the VLM segment.

Conditioning – Attention Sinks

Attention sinks → two separate Parameters

- `sink_pc` and `sink_action` (separate `nn.Parameter`s at width `d_pc` and `d_action`) sit at **position 0** of each segment, ahead of embodiment.
- Their job is to absorb softmax mass that heads can't usefully place – Xiao et al. 2023, applied per-expert.
- **Not shared**: PC bidirectional and Action causal have independent "good sink" directions; sharing weights would constrain a content-free Parameter for no gain.
- `sink_pc` also drains Action's cross-attention to PC K/V – double duty for free.

`AdaLNConditioner` is **timestep only** – embodiment and proprio-history adaLN paths are gone.

Token Sequence

Three segments with separate widths (`d_vlm` , `d_pc` , `d_action`), assembled by `SequenceBuilder` . Each token carries a type label in `{vlm, pc, action}` .

```
VLM segment (rho = vlm, width = d_vlm)
```

```
<im_start>system ... <im_end>
<im_start>user
  <vision_start> <image_pad>xN<vision_end>      cam 1 image
  <vision_start> <image_pad>xN<vision_end>      cam 2 image
  <vision_start> <image_pad>xN<vision_end>      cam 3 image
  Robot: {description}
  <task>{instruction}</task>
<im_end>
<im_start>assistant
```

```
PC segment (rho = pc, width = d_pc)
```

```
sink_pc          attention sink (Xiao et al. '23)
embodiment_pc    from SharedEmbodimentEmbedding
<align>
<pc_start> pc_cam1_tok_1 ... pc_cam1_tok_M <pc_end>      base cam
<pc_start> pc_cam2_tok_1 ... pc_cam2_tok_M <pc_end>      base cam
<pc_wrist_start> pc_cam3 ... pc_cam3_tok_M <pc_wrist_end> wrist
```

```
Action segment (rho = action, width = d_action)
```

```
sink_action          attention sink (Xiao et al. '23)
embodiment_action    from SharedEmbodimentEmbedding
<proprio_start> proprio_tok_1 ... proprio_tok_P <proprio_end>
<action_start>  act_tok_1 ... act_tok_T <action_end>
```


Token Sequence – Notes

Not in the token sequence – enter through adaLN only:

- Timestep embedding (sole adaLN signal)

For **K** cameras: VLM prompt has **K** image blocks and PC segment has **K** PC blocks – **one-to-one correspondence**.

Each camera is bracketed by `<pc_start>` / `<pc_end>` (non-wrist) **or** `<pc_wrist_start>` / `<pc_wrist_end>` (wrist), selected by a per-camera `is_wrist` flag passed to `SequenceBuilder`. Within each class the embeddings are **shared across cameras**, the same way Qwen-VL shares `<vision_start>` / `<vision_end>` across images. Wrist delimiters apply to the **PC segment only** – VLM-side wrist tokens are deferred to Phase 2 (VLM stays frozen).

Concrete Layout – 3 cameras (2 base + 1 wrist), 3 point clouds

What the Mixture-of-Transformers actually sees in one batch element.

```
VLM SEGMENT      rho = vlm      width d_vlm = 2560 (Qwen3-VL-4B-Instruct)
~700 tokens for K=3 cams (n_img ~196 merged ViT patches per image)

<im_start>system ...system prompt tokens... <im_end>
<im_start>user
  <vision_start> <image_pad> * n_img <vision_end>      ← base cam 1
  <vision_start> <image_pad> * n_img <vision_end>      ← base cam 2
  <vision_start> <image_pad> * n_img <vision_end>      ← wrist cam 3
  Robot: {description}
  <task> {instruction} </task>
<im_end>
<im_start>assistant

PC SEGMENT      rho = pc      width d_pc = 1024
length = 3 + K(M+2) = 3 + 3*(M+2) e.g. M=128 → 393 tokens

[0]      sink_pc      attention sink (learnable param)
[1]      embodiment_pc SharedEmbodimentEmbedding.to_pc(rid)
[2]      <align>      AlignmentToken(0|1)·d_pc
[3]      <pc_start>
[4..]    pc_cam1_tok_1 ... pc_cam1_tok_M ] cam 1 BASE (Uni3D output)
[.]      <pc_end>
[.]      <pc_start>
[..]    pc_cam2_tok_1 ... pc_cam2_tok_M ] cam 2 BASE
[.]      <pc_end>
[.]      <pc_wrist_start>
[..]    pc_cam3_tok_1 ... pc_cam3_tok_M ] cam 3 WRIST (separate delimiters)
[.]      <pc_wrist_end>

ACTION SEGMENT  rho = action width d_action = 1024
length = P + T + 6 e.g. P=1, T=30 → 37 tokens

[0]      sink_action      attention sink (learnable param)
[1]      embodiment_action SharedEmbodimentEmbedding.to_action(rid)
[2]      <proprio_start>
[3..]    proprio_tok_1 ... proprio_tok_P ProprioEncoder(state) ← variable P
[P+3]    <proprio_end>
[P+4]    <action_start>
[P+5..]  a_t ... a_{t+T-1}      noised actions a^τ = τ·a + (1-τ)·ε
[P+T+5]  <action_end>
```

Self-attention per segment: VLM causal (Qwen native); PC bidirectional; Action causal. **Cross-attention at expert layers:** PC reads VLM K/V (detached); Action reads VLM K/V (detached) + PC K/V (gradient flows). **adaLN** (DiT zero-init, 6 mod params) modulates the Action expert from the timestep only.

Special Tokens

Follows Qwen-VL's delimiter pattern: delimiter tokens wrap each segment, pad tokens are replaced by encoder features.

Qwen-VL images:	<code><vision_start></code>	<code><image_pad> xN</code>	<code><vision_end></code>	per camera
PC (non-wrist):	<code><pc_start></code>	<code><pc_pad> xM</code>	<code><pc_end></code>	per non-wrist cam
PC (wrist):	<code><pc_wrist_start></code>	<code><pc_pad> xM</code>	<code><pc_wrist_end></code>	per wrist cam
Proprio block:	<code><proprio_start></code>	<code>proprio_tokens</code>	<code><proprio_end></code>	wraps P proprio tokens (variable)
Our actions:	<code><action_start></code>	<code>action_tokens</code>	<code><action_end></code>	

Token	Expert	Shared across cams	Purpose
<code>sink_pc</code>	PC	one per sample	Learnable attention sink at PC position 0; drains softmax mass
<code>sink_action</code>	Action	one per sample	Learnable attention sink at Action position 0; drains softmax mass

Supported Embodiments

Robot	Description	DOF
Trossen	7-DOF Trossen arm with parallel gripper	7
Franka	7-DOF Franka Panda arm with parallel gripper	7
SO-101	6-DOF SO-101 arm with parallel gripper	6

Each description is inserted verbatim into the VLM prompt as `Robot: {description}`. A separate learned per-robot token from `SharedEmbodimentEmbedding(3, d_emb=32)` is spliced into the **PC and Action** streams (not the VLM segment) – at width `d_pc` and `d_action` respectively, both projected from one shared bottleneck.

VLM Prompt Template

`build_vlm_prompt` returns a single string. No embodiment splice on the VLM side – `SharedEmbodimentEmbedding` lives in the PC and Action streams.

```
<im_start>system
You are a robot manipulation assistant...<im_end>
<im_start>user
<vision_start> <image_pad> <vision_end>
<vision_start> <image_pad> <vision_end>
<vision_start> <image_pad> <vision_end>
Robot: {robot_description}
<task>{task_instruction}</task> <im_end>
<im_start>assistant
```

VLM Prompt – Notes

- **System prompt is identical** across all samples → cacheable across a batch.
- One `<vision_start> <image_pad> <vision_end>` set per camera (K sets total, matching K point clouds).
- Proprio is **no longer textualized** – it enters as Action-stream tokens via `ProprioEncoder`.
- Embodiment is **not in the VLM segment** – `embodiment_pc` and `embodiment_action` are spliced into the PC and Action segments via `SharedEmbodimentEmbedding`.
- All tokens from `<im_start>system` through `<im_start>assistant` carry `rho = vlm`.

Input Encoders

All encoders are **frozen**. Outputs are precomputed and cached by the dataloader. K cameras, each providing one image and one point cloud.

Input	Encoder	Projection
Images ($\times K$) + text	Qwen-VL ViT + tokenizer	native <code>d_vlm</code> (cached)
Point clouds ($\times K$, one per cam)	Pretrained PC encoder	<code>Linear(enc_dim, d_pc)</code> per cam + <code>CentroidPosEmbed(xyz $\rightarrow$ d_pc)</code> added
Noisy action chunk	– (raw, noised at <code>t</code>)	<code>Linear(action_dim, d_action)</code>
Proprio ($P \times 10$)	– (raw floats)	<code>ProprioEncoder</code> : <code>Linear(proprio_dim, d_action)</code> $\rightarrow$ action stream

`CentroidPosEmbed` mirrors Uni3D's input `pos_embed` (`Linear(3, 128)  $\rightarrow$  GELU  $\rightarrow$  Linear(128, d_pc)`); applied to cached FPS centroids and **added** to the projected Uni3D tokens before `SequenceBuilder`. Final `Linear` is zero-init $\rightarrow$ step-0 contribution is exactly zero.

Qwen2.5-VL-3B reference

```
d_vlm = 2048    heads = 16 Q / 2 KV (GQA ratio 8)    head_dim = 128
intermediate_size = 11008                            num_layers = 36
```

Expert Placement

Experts are placed via a single `expert_layer_ids` list – **both PC and Action experts go at the same layers**. If not specified, experts are placed at all layers.

```
create_backbone(  
    qwen_model,  
    d_pc=1024, d_action=1024, d_cond=64,  
    expert_layer_ids=[0, 4, 8, 12, 16, 20, 24, 28, 32, 35],  
    expert_mode="dense",    # or "sparse"  
)
```

Dense vs Sparse

- **Dense**: experts at every layer. Cross-attn to VLM at `expert_layer_ids`; self-attn only elsewhere. **Same trainable param count** as "all-cross."
- **Sparse**: experts only at `expert_layer_ids`. Pass-through at other layers. Trainable params **proportional** to number of expert layers.


Dense vs Sparse – Concrete Example

8 layers, experts at `[0, 4, 7]`:

Mode	Trainable	Sparse / All
All	342.5M	100 %
Dense	342.5M	100 %
Sparse	128.4M	37.5 %

"All" and "Dense" have identical trainable counts – the difference is only in the attention pattern at non-expert layers (cross vs self-only). Sparse skips the experts entirely at non-expert layers.

Architecture Summary



Architecture Overview

Cached Encoders (precomputed by dataloader, K cameras)



Project to expert widths (d_{vlm} , d_{pc} per camera, d_{action})



Backbone (L layers, VLM runs at every layer)

Expert layer ($\text{is_cross}=\text{True}$):

VLM: Qwen decoder layer (frozen) → extract K/V from cache
PC: self-attn + cross-attn to VLM K/V → produces pc_k
Action: self-attn(causal) + cross-attn to [VLM K/V, pc_k]
+ adaLN (timestep only)

Non-expert layer (dense, $\text{is_cross}=\text{False}$):

VLM: Qwen decoder layer (frozen)
PC: self-attn only (no VLM K/V)
Action: self-attn only (causal, no VLM/PC K/V)

Non-expert layer (sparse):

VLM: Qwen decoder layer (frozen)
PC: pass-through
Action: pass-through

Gradient isolation: VLM K/V .detach(). PC K/V flow grads.



PC output
(MSE loss)

Action output
(flow matching loss)

Cross-Attention Flow at Expert Layers

Both PC and Action experts sit at the same `expert_layer_ids`. At each expert layer, **PC runs first**, then **Action** uses that layer's VLM K/V **and** PC K/V:

```
expert_layer_ids = [0, 4, 8]
```

```
Layer 0 (expert):  PC → [vlm_k_0]          Action → [vlm_k_0, pc_k_0]
```

```
Layer 1-3:        (dense: self-attn only / sparse: pass-through)
```

```
Layer 4 (expert):  PC → [vlm_k_4]          Action → [vlm_k_4, pc_k_4]
```

```
Layer 5-7:        (dense: self-attn only / sparse: pass-through)
```

```
Layer 8 (expert):  PC → [vlm_k_8]          Action → [vlm_k_8, pc_k_8]
```

At each expert layer: **VLM runs → PC cross-attends to VLM K/V (detached) → Action cross-attends to VLM K/V (detached) + PC K/V (gradients flow)**. All three use the same layer's features.

Per-Layer Expert Details

Feature	VLM (Qwen layer)	PC (ExpertBlock)	Action (ExpertBlock)
Norm	Qwen2RMSNorm	Qwen2RMSNorm	Qwen2RMSNorm
RoPE	3D MRoPE (native)	MRoPE chunk-shared (cont. from VLM)	MRoPE 1D (cont. from PC)
GQA	native (16Q/2KV for 3B)	proportional to width	proportional to width
MLP	GatedMLP (native)	GatedMLP (proportional)	GatedMLP (proportional)
Self-attn	Causal + packed-seq	Bidirectional	Causal
Cross-attn	N/A	VLM K/V detached (expert layers)	VLM K/V detached + PC K/V w/ grad (expert layers)
adaLN	None	None	timestep only
Status	Frozen	Trainable (init from VLM)	Trainable (init from VLM)

Position-ID Rule

A **single global counter** p walks $\text{VLM} \rightarrow \text{PC} \rightarrow \text{Action}$. Each token type advances the counter differently:

Token	(t, h, w)	Counter
VLM text	(p, p, p)	$p += 1$
VLM image patch (row r , col c of $H \times W$)	$(p, p+r, p+c)$	$p += \max(H, W)$ once after the whole image
All delimiters ($\langle \text{pc_start/end} \rangle$, $\langle \text{pc_wrist_start/end} \rangle$, $\langle \text{align} \rangle$, $\langle \text{proprio_start/end} \rangle$, $\langle \text{action_start/end} \rangle$, ...)	(p, p, p)	$p += 1$
Per-segment scalars (sink_pc , embodiment_pc , sink_action , embodiment_action)	(p, p, p)	$p += 1$
PC chunk (M tokens from one Uni3D run)	all M share (p, p, p)	$p += 1$ once for the whole chunk
Proprio token	(p, p, p)	$p += 1$
Action token	(p, p, p)	$p += 1$

Since $t = h = w$ everywhere except VLM image patches, MRoPE collapses to **1-D RoPE** on the PC and Action segments.

PC-Chunk MRoPE – Why Chunk-Shared

Inside one PC chunk, all M tokens carry **identical** $(p, p, p) \rightarrow$ pairwise $\Delta p = 0 \rightarrow$ MRoPE contributes **no positional signal within a chunk**. That's intentional – intra-chunk geometry is supplied by **two channels**: (1) Uni3D's pretrained `pos_embed` + 24 ViT blocks already baked it into the cached patch tokens, and (2) `CentroidPosEmbed` re-injects an explicit, fresh signal at PC-expert entry. **Across** chunks, $\Delta p \geq 3$ (separated by `<pc*_end>` and `<pc*_start>`) $\rightarrow$ MRoPE *does* carry the relative-position signal, so cross-camera attention can tell chunks apart and message-pass between them.

seq idx	token	t	h	w	
0	sink_pc	0	0	0	
1	embodiment_pc	1	1	1	
2	<align>	2	2	2	
3	<pc_start>	3	3	3	
4..	pc1_tok_1..M	4	4	4	← chunk 1: all M share (4,4,4); counter +1 once
M+4	<pc_end>	5	5	5	
M+5	<pc_start>	6	6	6	
M+6..	pc2_tok_1..M	7	7	7	← chunk 2: all M share (7,7,7); counter +1 once
2M+6	<pc_end>	8	8	8	

Two pretrained Uni3D encoders (base cams, wrist cams) – tokens come from the dataloader at width `d_pc`.

PC-Chunk MRoPE – API

`build_pc_chunk_position_ids(pc_chunk_sizes, start)` in `core.py` returns the 1-D position list plus `next_p` for the action start:

```
positions, next_p = build_pc_chunk_position_ids([4, 4], start=10)
# positions = [10, 11, 12,          sink, emb, align
#             13, 14, 14, 14, 14, 15,    <s>, [4]×4, <e>
#             16, 17, 17, 17, 17, 18]    <s>, [4]×4, <e>
# next_p     = 19                      (action segment starts here)
```

`Backbone.forward(..., pc_chunk_sizes=[M1, ..., MK])` activates the rule. `pc_chunk_sizes=None` (default) preserves the legacy per-token 1-D continuation.

Layout (matches `SequenceBuilder.forward`): `sink_pc` → `emb_pc` → `<align>` each `p+=1`, then per camera `<pc*_start>` `p+=1`, **M chunk tokens share p** then `p+=1`, `<pc*_end>` `p+=1`.

CentroidPosEmbed – Re-injecting Intra-Chunk Geometry

Chunk-shared MRoPE is **silent within a PC chunk**. Intra-chunk geometry then depends on Uni3D's `pos_embed` surviving 24 ViT blocks – thin. `CentroidPosEmbed` adds a **fresh, explicit** signal at PC-expert entry, mirroring Uni3D's own recipe:

```
class CentroidPosEmbed(nn.Module):          # core.py
    def __init__(self, d_pc: int, hidden: int = 128):
        self.fc1 = nn.Linear(3, hidden)
        self.act = nn.GELU()
        self.fc2 = nn.Linear(hidden, d_pc)
        nn.init.zeros_(self.fc2.weight)      # adaLN-zero recipe
        nn.init.zeros_(self.fc2.bias)         # step-0 output  $\equiv 0$ 

    def forward(self, centroids):             # (B, M, 3)  $\rightarrow$  (B, M, d_pc)
        return self.fc2(self.act(self.fc1(centroids)))
```

Per-camera composition before `SequenceBuilder`:

```
pc_tokens_per_cam[k] = pc_proj(features[k]) + centroid_pos_embed(centroids[k])
```

- **Zero-init final Linear** $\rightarrow$ step-0 output $\equiv 0 \rightarrow$ projected PC tokens are bit-for-bit identical to the no-centroid baseline (preserves the VLM-init prior on PC ExpertBlocks).
- **Hidden = 128** mirrors Uni3D's `Linear(3, 128)  $\rightarrow$  GELU  $\rightarrow$  Linear(128, embed_dim)` recipe; **shared across cameras** (wrist + base).

adaLN Conditioning – DiT adaLN-zero (6 mod params)

Only on the **Action** expert, and now only **timestep** flows in. Embodiment lives in the PC + Action streams (**SharedEmbodimentEmbedding**); proprio enters as Action-stream tokens.

Per-layer MLP produces **6** mod params per ExpertBlock – DiT adaLN-zero (Peebles & Xie 2023):

```
# Conditioning vector
cond = TimestepEmbedding(t)          # (B, d_t)

# Per-layer modulation (inside each MultiExpertLayer)
s1, sh1, g1, s2, sh2, g2 = action_adaln(cond).chunk(6)  # each (B, d_action)

# Attention sub-block (gated residual)
normed = (1 + s1) * RMSNorm(x) + sh1
attn_out = o_proj(attention(normed, ...))
x = x + g1 * attn_out

# MLP sub-block (gated residual)
normed2 = (1 + s2) * RMSNorm(x) + sh2
mlp_out = mlp(normed2)
x = x + g2 * mlp_out
```

Attention Masks

Query \ Key	VLM	PC	Action
VLM	causal + packed-seq	blocked	blocked
PC	full (detached K/V)	bidirectional	blocked
Action	full (detached K/V)	full (PC K/V w/ grad)	causal

- Action uses a **hybrid mask** at expert layers: full cross to VLM+PC columns, causal on self columns.
- At non-expert layers (dense), Action uses **causal self-only** mask.
- Head count mismatches handled by `_match_heads` (GQA-style grouping).

Gradient Isolation

Flow matching loss	→ Action expert + adaLN + PC expert (via PC K/V)	--X VLM (detached)
MSE loss	→ PC expert	--X Action, VLM (detached)

Detach points:

- VLM K/V from `DynamicCache` (read-only for experts)
- VLM params frozen end-to-end

PC K/V are **not** detached when Action consumes them – action flow-matching loss now flows back into the PC expert. The PC expert receives gradient from both losses (its own MSE and the action loss). MSE doesn't reach Action because PC's forward never consumes Action K/V.

Trainable vs Frozen

Frozen

- Qwen decoder layers (all L)
- Final RMSNorm
- Rotary embeddings
- ViT (dataloader)
- PC encoder (dataloader)

Trainable

- PC / Action `ExpertBlock` s (selected layers)
- Action adaLN per layer (timestep only)
- Input projections (incl. `ProprioEncoder`)
- Special token embeddings
- Alignment embedding
- `SharedEmbodimentEmbedding`: `base:`
`Embedding(3, d_emb)` + zero-init `Linear`
heads to `d_pc` & `d_action`
- Per-expert attention sinks: `sink_pc` ,
`sink_action` (separate `nn.Parameter` s)

Param Count Example

Qwen2.5-VL-3B, 2 layers, `d_pc = d_action = 1024`, all-cross:

VLM (frozen):	154.15 M
PC (trainable):	38.54 M
Action + adaLN:	47.07 M
Total trainable:	85.61 M (35.7 %)

This is a minimal 2-layer config; at full 36 layers the trainable share stays similar since PC/Action widths scale together with depth.

Training

Dataloader provides: cached VLM embeddings (text + vision merged), PC features, action ground truth.

1. **Project** PC/Action to expert widths; **noise** action at random timestep **t**.
2. **Prepare** VLM inputs (position IDs, attn mask, RoPE) and expert positions (MRoPE continuation).
3. **Forward** through backbone – VLM at every layer, experts at configured layers.
4. **Losses**: flow matching on Action output, MSE on PC output. **No loss on VLM.**

Inference

1. **Full forward** through all layers – VLM + PC computed once, K/V cached.
2. **T denoising steps**: only action tokens re-enter, attending to cached VLM/PC K/V.
3. **adaLN** injects current timestep at each step.

VLM and PC compute exactly once per observation. Only the Action expert runs T times during denoising, making inference cost dominated by action-chunk length and T .

Design Rationale – Part 1

Actual Qwen layers for VLM

Zero distribution shift, bit-for-bit verified identical output.

RMSNorm + MRoPE on experts

Matches VLM primitives – initialized weights are compatible from step 0.

PC bidirectional, Action causal

Point clouds have no ordering; action chunks are temporal sequences.

Design Rationale – Part 2

Dense vs sparse modes

Dense gives full-depth processing with selective cross-attention; Sparse minimizes trainable params for faster iteration.

Cached encoder outputs

ViT and PC encoder are frozen – run once per sample in preprocessing.

Design Rationale – Part 3

Proprio as Action-stream tokens

A single `Linear(proprio_dim, d_action)` projects each proprio reading and the result is **prepended** to the action chunk under a unified causal mask. Same Linear handles current-only or current+history without architectural change. Action tokens can attend back to proprio while staying causal among themselves.

Embodiment as a shared bottleneck across PC + Action

`Robot: {description}` text gives the **frozen VLM** semantic context. A separate `SharedEmbodimentEmbedding` – one `nn.Embedding(3, d_emb)` base + zero-init `Linear` heads – splices `embodiment_pc` and `embodiment_action` at position 0 of the PC and Action segments. Both losses train the shared base; Action loss also reaches `to_pc` via the no-detach PC→Action cross-attention. Embodiment is deliberately **not** in the VLM segment (gradient cannot reach into a frozen stack with detached cross-attention K/V).

Alignment flag

Binary `Embedding` (not a continuous projection) because the signal **is discrete** – point clouds are either registered or not.

Code Structure

```
unified_vla/  
├── core.py      TokenType, TokenSequence, SequenceBuilder, AlignmentToken,  
                InputProjections, SharedEmbodimentEmbedding, ProprioEncoder,  
                AdaLNConditioner (timestep-only), CentroidPosEmbed, modulate  
├── attention.py ExpertQKV (GQA), build_attention_mask, _match_heads,  
                detached_cross_modal_attention  
├── layers.py    GatedMLP, ExpertBlock (RMSNorm + MRoPE + mask + adaLN),  
                init_expert_from_vlm  
├── backbone.py  MultiExpertLayer, Backbone (dense/sparse, expert_layer_ids)  
├── surgery.py   create_backbone(qwen_model, ... ),  
                init_special_tokens_from_vlm  
├── losses.py    TimestepEmbedding, noise_action,  
                flow_matching_loss, pc_mse_loss  
├── prompt.py    build_vlm_prompt (single string)  
└── utils.py     count_params, backbone_param_summary, print_param_summary
```

Open Questions

- **Expert width sweep:** `d_pc` and `d_action` from 512 → 2048, asymmetric configs.
- **Expert layer placement:** which layers matter most for cross-attention? Every 2nd? Every 4th? Last N?
- `flex_attention`: replace Action hybrid bool mask with `torch.nn.attention.flex_attention` for efficiency.
- **Phase-2 VLM unfreezing:** learning-rate ratio, what to do with VLM text output.
- **Fused attention:** single call with segment-level masking vs three separate calls.

Implementation Status

243 tests (unit + integration), all passing.



Completed

- ✓ Token type registry, sequence builder (proprio prefix), alignment token (`core.py`)
- ✓ Input projections, `ProprioEncoder` , `SharedEmbodimentEmbedding` , timestep-only `AdaLNConditioner` (`core.py`)
- ✓ `modulate` function (`core.py`)
- ✓ `ExpertQKV` with GQA, attention mask builder, gradient-detached cross-modal attention (`attention.py`)
- ✓ `ExpertBlock` with RMSNorm + MRoPE + GatedMLP + adaLN + mask support (`layers.py`)
- ✓ Expert weight initialization from VLM, GQA-aware slicing (`layers.py`)
- ✓ `MultiExpertLayer` : actual Qwen layer + `ExpertBlock` s, dense/sparse modes (`backbone.py`)
- ✓ `Backbone` with `expert_layer_ids` , dense/sparse modes (`backbone.py`)
- ✓ `create_backbone` surgery from Qwen2.5-VL + `init_special_tokens_from_vlm` (`surgery.py`)
- ✓ VLM prompt builder (`prompt.py`) – single string (no embodiment splice on VLM side)
- ✓ Flow matching loss, PC MSE loss, timestep embedding (`losses.py`)
- ✓ Param-counting utilities (`utils.py`)
- ✓ VLM forward exactness: bit-for-bit identical to Qwen2.5-VL-3B at every layer
- ✓ Dense/sparse mode tests with param-count verification
- ✓ **Proprio as Action-stream tokens** – `ProprioEncoder` prepends projected proprio to the action chunk under a unified causal mask; adaLN shrunk to timestep-only

Remaining – Pipeline Completion

- ■ Combined training step (forward + both losses + optimizer)
- ■ Inference with K/V caching (VLM/PC cached, action iterates)
- ■ End-to-end smoke test

Remaining – Architecture Backlog (Phase 1)

- **Wrist-camera Uni3D encoder** – second pretrained Uni3D for wrist clouds (close-up, gripper-scale geometry); tokens come from dataloader; backbone treats wrist chunks identically (same `pc_chunk_sizes` rule, separate `<pc_wrist_*` delimiters already wired).
- **Paired camera shuffle** – DataLoader aug; same permutation on image + PC streams (and the matching `is_wrist_per_camera` list).
- **Paired camera dropping** – precompute `E_zero = ViT(zero_image)` once; swap dropped cameras' cached embedding with `E_zero`; keep ≥ 1 camera.
- **Finetune PC encoder flag** – `finetune_pc_encoder: bool = False`; disables caching when enabled. Direction to explore.

Thank You

Questions?

Unified VLA · Single-Backbone Multi-Expert Architecture Qwen2.5-VL + PC expert + Action flow-matching expert

Appendix A

MRoPE Deep Dive

How multimodal rotary position embedding works,
from 1-D RoPE up to three-axis spatial assignment.

MRoPE at a Glance

A **single position counter** p stamps every token in the prompt.

- **Text token** $\rightarrow (t, h, w) = (p, p, p)$, then $p \leftarrow p + 1$
- **Image of grid** $H \times W$ (still image): all patches share $t = p$; patch at row r , col c gets $(t, h, w) = (p, p + r, p + c)$; counter jumps $p \leftarrow p + \max(H, W)$
- `<vision_start>` / `<vision_end>` are regular **text** tokens – they follow the text rule
- Inside each attention head, the head_dim pairs are **partitioned across three axes** via `mrope_section`.
Each pair has its own axis and its own frequency

Frequency Spectrum

Every pair draws its frequency from one geometric spectrum (same as 1-D RoPE):

$$\theta_i = 10000^{-2i/d}, \quad i \in \{0, 1, \dots, d/2 - 1\}$$

For Qwen2.5-VL-3B, `d = head_dim = 128`, so `i ∈ [0, 63]`:

i	θ_i	Wavelength $2\pi/\theta_i$	Axis (see next slide)
0	1.000	6.28	t
1	0.867	7.25	t
16	0.100	63	h
32	0.010	628	h
63	1.2×10^{-4}	51,500	w

Fast pairs resolve local offsets; slow pairs resolve long-range ones.

Head-Dim Accounting – Qwen2.5-VL-3B

```
head_dim      = 128
number of pairs = 128 / 2 = 64
mrope_section  = [16, 24, 24]
```

Pair index → axis assignment:

```
pair i  ∈ [ 0 ..... 15 | 16 ..... 39 | 40 ..... 63 ]
axis    ∈ [   t      |   h      |   w      ]
count   ∈ [   16     |   24     |   24     ]
```

Each pair is fully specified by its index **i**:

$$\text{axis}(i) = \begin{cases} t & i \in [0, 15] \\ h & i \in [16, 39] \\ w & i \in [40, 63] \end{cases} \quad \theta_i = 10000^{-2i/128}$$

1-D RoPE – Rotation Formula

Treat each Q/K vector as a stack of **2-D pairs**. For position m and frequency θ_i , rotate each pair by angle $m \cdot \theta_i$:

$$R(\alpha) \cdot (x, y) = (x \cos \alpha - y \sin \alpha, x \sin \alpha + y \cos \alpha)$$

In components, for pair (q_{2i}, q_{2i+1}) at position m :

$$\begin{aligned} q'_{2i} &= q_{2i} \cos(m\theta_i) - q_{2i+1} \sin(m\theta_i) \\ q'_{2i+1} &= q_{2i} \sin(m\theta_i) + q_{2i+1} \cos(m\theta_i) \end{aligned}$$

The Relativity Identity

After rotating q at position m and k at position n , the dot product depends **only on the offset** $n - m$:

$$\langle R(m\theta)q, R(n\theta)k \rangle = \langle q, R((n - m)\theta)k \rangle$$

Why this works. Rotations are orthogonal, so

$$R(m\theta)^\top R(n\theta) = R(-m\theta) R(n\theta) = R((n - m)\theta)$$

For $q = k = (1, 0)$ (unit pair), the dot product collapses even further to a scalar cosine:

$$\langle R(m\theta)q, R(n\theta)k \rangle = \cos((n - m)\theta)$$

This is the whole "relative position via rotation" trick. MRoPE just applies it independently on three axes.

1-D RoPE in PyTorch (Qwen / LLaMA style)

The HF implementation pairs dim `i` with dim `i + d/2` (not adjacent) via `rotate_half`:

```
def rotate_half(x: torch.Tensor) → torch.Tensor:
    """Splits the last dim in half and rotates: [a | b] → [-b | a]."""
    half = x.shape[-1] // 2
    x1, x2 = x[ ..., :half], x[ ..., half:]
    return torch.cat((-x2, x1), dim=-1)

def apply_rotary_pos_emb(
    q: torch.Tensor,      # (B, H, T, D)
    k: torch.Tensor,      # (B, H_kv, T, D)
    cos: torch.Tensor,    # (B, T, D) – precomputed cos(m·θi), mirrored across halves
    sin: torch.Tensor,    # (B, T, D) – precomputed sin(m·θi), mirrored across halves
):
    cos = cos.unsqueeze(1) # broadcast over heads
    sin = sin.unsqueeze(1)
    q_rot = (q * cos) + (rotate_half(q) * sin)
    k_rot = (k * cos) + (rotate_half(k) * sin)
    return q_rot, k_rot
```

Mathematically equivalent to the "adjacent pair" formulation – just a reshuffled dim layout so the rotation is one elementwise multiply.

MRoPE – The Axis Split

Starting from the 1-D RoPE cos/sin tables (one per axis: $\cos[t]$, $\cos[h]$, $\cos[w]$), MRoPE **slices the head_dim by mrope_section and picks a different axis's table per slice**:

head_dim layout after split:

first half (dims 0 ... 63)			
t (16)		h (24)	w (24)
second half (dims 64 ... 127)			
t (16)		h (24)	w (24)

The **second half repeats** the same $[t, h, w]$ pattern because **rotate_half** pairs dim i with dim $i + 64$ – both halves of one pair must share an axis.

Score between tokens **A** and **B** decomposes by axis:

$$\text{score}(A, B) = \underbrace{\sum_{i \in t\text{-block}} \cos(\Delta t \cdot \theta_i)}_{16 \text{ pairs}} + \underbrace{\sum_{i \in h\text{-block}} \cos(\Delta h \cdot \theta_i)}_{24 \text{ pairs}} + \underbrace{\sum_{i \in w\text{-block}} \cos(\Delta w \cdot \theta_i)}_{24 \text{ pairs}}$$

MRoPE in PyTorch (from `modeling_qwen2_5_vl.py`)

```
def apply_multimodal_rotary_pos_emb(q, k, cos, sin, mrope_section, unsqueeze_dim=1):
    """
    cos, sin have shape (3, B, T, head_dim) – axis 0 indexes (t, h, w).
    mrope_section e.g. [16, 24, 24] for Qwen2.5-VL-3B.
    """
    # Duplicate [16, 24, 24] → [16, 24, 24, 16, 24, 24] for rotate_half's two halves
    mrope_section = mrope_section * 2

    # Split cos / sin along head_dim and pick axis i % 3 for each chunk
    cos = torch.cat(
        [m[i % 3] for i, m in enumerate(cos.split(mrope_section, dim=-1))],
        dim=-1,
    ).unsqueeze(unsqueeze_dim)
    sin = torch.cat(
        [m[i % 3] for i, m in enumerate(sin.split(mrope_section, dim=-1))],
        dim=-1,
    ).unsqueeze(unsqueeze_dim)

    # Standard rotate-half application, now using axis-mixed cos/sin
    q_embed = (q * cos) + (rotate_half(q) * sin)
    k_embed = (k * cos) + (rotate_half(k) * sin)
    return q_embed, k_embed
```

The two interesting lines are the `cat + split + i % 3` pattern – that's **the entire axis-mixing logic**. Everything else is ordinary RoPE.

Position Assignment – The Counter Rule

Maintain a single scalar counter p , initialized to 0 . Walk the token stream:

```
for token in prompt:
    if token.is_text:
        assign (t, h, w) = (p, p, p)
        p += 1

    elif token.is_image:          # H x W patch grid
        for r in range(H):
            for c in range(W):
                assign (t, h, w) = (p, p + r, p + c)
            p += max(H, W)        # jump, not H*W
```

Three consequences:

- All patches of one still image share $t = p \rightarrow t$ becomes an "image ID"
- Different images get different t values $\rightarrow$ cross-image attention sees $\Delta t \neq 0$
- Text gap between the last image patch and the next text token is $\max(H, W)$, not $H \cdot W \rightarrow$ image density doesn't inflate long-range scale

Position Assignment – PyTorch Sketch

A simplified version of Qwen's `get_rope_index` for one image block:

```
def assign_image_positions(p_start: int, H: int, W: int) → torch.Tensor:
    """
    Returns position_ids of shape (3, H*W) for one still image starting at p_start.
    All patches share t = p_start; h and w vary by row and column.
    """
    t = torch.full((H * W,), p_start, dtype=torch.long)
    h = p_start + torch.arange(H).repeat_interleave(W) # 0,0,0,1,1,1,2,2,2,...
    w = p_start + torch.arange(W).repeat(H)             # 0,1,2,0,1,2,0,1,2,...
    return torch.stack([t, h, w])                       # (3, H*W)

def next_counter(p_start: int, H: int, W: int) → int:
    """Counter value for the token immediately after the image."""
    return p_start + max(H, W)
```

Full implementation (handling text interleaving, multiple images, videos with `second_per_grid_t`) lives in `Qwen2_5_VLModel.get_rope_index` (`modeling_qwen2_5_vl.py:956`).

Scenario A – Pure Text

Sequence: The cat sat

seq idx	token	t	h	w
0	The	0	0	0
1	cat	1	1	1
2	sat	2	2	2

For any two tokens at positions i , j : all three deltas are equal, $\Delta t = \Delta h = \Delta w = j - i$. The axis decomposition collapses:

$$\text{score}(i, j) = \sum_{k \in t} \cos((j-i)\theta_k) + \sum_{k \in h} \cos((j-i)\theta_k) + \sum_{k \in w} \cos((j-i)\theta_k) = \sum_{k=0}^{63} \cos((j-i)\theta_k)$$

Punchline. When $t = h = w$, MRoPE **degenerates to ordinary 1-D RoPE**. Text is a strict special case; the axis split buys you nothing.

Scenario B – One Image (Setup)

Sequence: See <vision_start> [6 patches, H=2 W=3] <vision_end> it.

seq idx	token	t	h	w
0	See	0	0	0
1	<vision_start>	1	1	1
2	patch (0, 0)	2	2	2
3	patch (0, 1)	2	2	3
4	patch (0, 2)	2	2	4
5	patch (1, 0)	2	3	2
6	patch (1, 1)	2	3	3
7	patch (1, 2)	2	3	4
8	<vision_end>	5	5	5
9	it	6	6	6
10	.	7	7	7

Image's max coord is $2 + \max(2, 3) - 1 = 4$, so <vision_end> starts at 5 – the counter jumped by 3, not 6.

Scenario B – What Attention Sees

Intra-image (e.g. patch $(2,2,2) \leftrightarrow \text{patch } (2,3,4)$):

$$\Delta t = 0, \quad \Delta h = 1, \quad \Delta w = 2$$

The 16 **t**-pairs all sit at $\cos(0) = 1$ – constant score, no information. The **24 h-pairs** and **24 w-pairs** carry the full 2-D spatial signal. Intra-image attention behaves like a pure 2-D positional prior.

Image $\leftrightarrow$ text (patch $(2,3,4) \leftrightarrow \text{it at } (6,6,6)$):

$$\Delta t = 4, \quad \Delta h = 3, \quad \Delta w = 2$$

All three blocks contribute – the model reads a mixed positional signature placing the image patch relative to the text position.

Text $\leftrightarrow$ text across the image (**See** $(0,0,0) \leftrightarrow \text{it } (6,6,6)$):

$$\Delta t = \Delta h = \Delta w = 6 \quad (\text{not } 10)$$

Scenario C – Two Images (First Image)

Sequence: Left <vs> [img1 2×3] <ve> vs <vs> [img2 2×3] <ve> right.

seq idx	token	t	h	w
0	Left	0	0	0
1	<vs>	1	1	1
2	img1 (0, 0)	2	2	2
3	img1 (0, 1)	2	2	3
4	img1 (0, 2)	2	2	4
5	img1 (1, 0)	2	3	2
6	img1 (1, 1)	2	3	3
7	img1 (1, 2)	2	3	4
8	<ve>	5	5	5
9	vs	6	6	6

After image 1: counter jumps 2 → 5. Text resumes at 5, 6, then <vs> of image 2 at 7.

Scenario C – Two Images (Second Image)

seq idx	token	t	h	w
10	<vs>	7	7	7
11	img2 (0, 0)	8	8	8
12	img2 (0, 1)	8	8	9
13	img2 (0, 2)	8	8	10
14	img2 (1, 0)	8	9	8
15	img2 (1, 1)	8	9	9
16	img2 (1, 2)	8	9	10
17	<ve>	11	11	11
18	right	12	12	12
19	.	13	13	13

Image 1 has **t = 2** ; **image 2** has **t = 8** . Different **t** values mark them as different images. Inside each image, the relative **(h, w)** structure is identical – translation-invariant across images.

Scenario C – What Attention Sees

Same corner, different images (img1 top-left (2,2,2) ↔ img2 top-left (8,8,8)):

$$\Delta t = 6, \quad \Delta h = 6, \quad \Delta w = 6$$

All three blocks agree on a large offset – the 16 **t**-pairs decisively separate the two images. **The **t**-axis acts as an image-id signal.**

Intra-image 2 diagonal (img2 (8,8,8) ↔ (8,9,10)):

$$\Delta t = 0, \quad \Delta h = 1, \quad \Delta w = 2$$

Identical to the intra-image 1 diagonal – the per-image spatial attention is translation-invariant. The model learns relational patterns like "bottom-right corner" without tying them to absolute coordinates.

Cross-image same-corner pair: (2,2,2) ↔ (8,9,10) → $\Delta t = 6, \Delta h = 7, \Delta w = 8$. Large on all three axes – "far apart in every sense."

Summary – The Three-Line Pattern

Scenario	Within one "block"	Across blocks
Pure text	All axes move together – MRoPE collapses to 1-D RoPE	N/A
One image	$\Delta t = 0$ inside image; h/w carry 2-D spatial signal	Text $\leftrightarrow$ image mixes all three axes; text $\leftrightarrow$ text compresses by $\max(H, W)$
Two images	Same as one image, repeated	Different images $\Rightarrow$ non-zero $\Delta t \Rightarrow t$ -axis labels "which image"; relative (h, w) is translation-invariant

Mental hook. The t -axis is an image-identity / time axis. The (h, w) axes carry 2-D spatial layout inside each image. When a token is text, all three collapse to one sequential axis. The $\max(H, W)$ jump keeps image density from inflating text's long-range position scale. Everything else is **the 1-D RoPE machinery from earlier, run three times in parallel** with the pair-to-axis mapping from `mrope_section`.

Appendix B

Uni3D – Point Cloud Encoder

How a standard 2-D ViT is re-used for 3-D point clouds,
with emphasis on the positional encoding design.

Xue et al., ICLR 2024 · arXiv 2310.06773 · [baaivision/Uni3D](#)

The Big Idea

Problem. 3-D foundation models start from scratch – tiny datasets (~1 M shapes), no pretrained-ViT inheritance.

Uni3D's trick. Turn a point cloud into a **sequence of tokens** so a standard 2-D-pretrained ViT (e.g. EVA / CLIP-ViT) can process it **unchanged** – then align its **CLS** output to CLIP image-text features.

- **Input:** **xyz** coords $(B, N, 3)$ + per-point **colors** $(B, N, 3)$ with $N \approx 10\,000$ points.
- **Output:** one embedding per point cloud $(B, 512)$ aligned to CLIP text/image space.
- **Training target:** CLIP-style contrastive on triplets (point cloud, image, text).
- **Scale:** up to **~1 B params** on ~10 M images / ~70 M texts / ~1 M shapes.

The "3-D-ness" is confined to **two small, non-ViT pieces**: a PointNet patchifier and a positional MLP on centroids. Everything else is a stock ViT loaded straight from **timm**.

Pipeline at a Glance

Four stages – input to CLIP-aligned embedding:

xyz : (B, N, 3) N ≈ 10 000 points
color : (B, N, 3)

1. Group (FPS + kNN)
centers (B, 512, 3)
neighborhoods (B, 512, 32, 3) centered per-group
features (B, 512, 32, 6) xyz + rgb

2. Encoder (residual PointNet)
patch tokens (B, 512, 256) → Linear 256→768

3. Add positional encoding + prepend CLS
pos = MLP(centers) (B, 512, 768)
x = CAT(cls, tokens) + CAT(cls_pos, pos)

4. ViT blocks (reused from 2-D) → CLS → Linear 768→512

embedding : (B, 512) CLIP-aligned

Stage 1 – Group (FPS + kNN)

```
class Group(nn.Module):
    def __init__(self, num_group=512, group_size=32):
        ...

    def forward(self, xyz, color):
        # Farthest-point sampling picks num_group = 512 coverage centers
        center = fps(xyz, self.num_group)           # (B, 512, 3)

        # k-nearest-neighbors gathers group_size = 32 points per center
        idx = knn_point(self.group_size, xyz, center) # (B, 512, 32)
        neighborhood = gather(xyz, idx)              # (B, 512, 32, 3)
        neighborhood_color = gather(color, idx)      # (B, 512, 32, 3)

        # CRITICAL: subtract the centroid so each patch is translation-invariant
        neighborhood = neighborhood - center.unsqueeze(2)

        features = torch.cat([neighborhood, neighborhood_color], dim=-1) # (B, 512, 32, 6)
        return neighborhood, center, features
```

Why this matters.

- **FPS** spreads the 512 centers evenly – no hyperparameter-sensitive voxelization.
- **kNN** gives a fixed-size local context per center → every patch has the same shape.
- **Centroid subtraction** decouples **shape** (stored in patch content) from **location** (stored only in **center**, which later feeds the positional MLP). Analogous to how an image patch's pixels don't know where they are in the image.

Stage 2 – Patch Embedding (Residual PointNet)

A mini-PointNet that turns (32, 6) local coordinates into one 256-D token per patch:

```
class Encoder(nn.Module):
    def __init__(self, encoder_channel=256):
        self.first_conv = Conv1d(6, 128) → BN → ReLU → Conv1d(128, 256)
        self.second_conv = Conv1d(512, 512) → BN → ReLU → Conv1d(512, 256)

    def forward(self, point_groups):
        # (B·G, 32, 6)
        x = self.first_conv(x.T) # (B·G, 256, 32)
        g = x.max(dim=2, keepdim=True) # (B·G, 256, 1) ← global per-patch feature
        x = cat([g.expand(-1, -1, 32), x]) # (B·G, 512, 32) ← residual-style concat
        x = self.second_conv(x) # (B·G, 256, 32)
        return x.max(dim=2) # (B·G, 256) ← final per-patch token
```

The residual-max-pool trick (PointBERT-style).

1. First Conv1d lifts each point to 256-D.
2. **Max-pool gives a global patch descriptor** (permutation-invariant over points).
3. **Concat global with per-point** so the second Conv1d sees "how does this point differ from the patch summary?"
4. Second max-pool → one permutation-invariant token per patch.

Result: (B, 512, 256), then Linear(256 → 768) lifts to ViT width.

Stage 3 – Positional Encoding (the interesting part)

```
self.pos_embed = nn.Sequential(
    nn.Linear(3, 128),
    nn.GELU(),
    nn.Linear(128, 768),      # trans_dim = ViT width
)

pos = self.pos_embed(center) # (B, 512, 3) → (B, 512, 768)
```

A **learned 2-layer MLP** applied to each patch's **3-D centroid coordinate**.

- **Not sinusoidal** – no $\sin(\omega x) / \cos(\omega x)$ table.
- **Not a lookup** – image ViT uses `nn.Embedding(num_patches, d)` indexed by patch position; here positions are **continuous** $(x, y, z) \in \mathbb{R}^3$, so we use a function $\mathbb{R}^3 \rightarrow \mathbb{R}^{768}$.
- **Absolute position lives only here**; patch content is translation-invariant (Stage 1's centroid subtraction).
- Continuous MLP → generalizes to shapes of different spatial extent without retraining the pos table.

Why an MLP for 3-D Position?

	Image ViT	Uni3D point cloud
Grid structure	Fixed $H \times W$ grid	No grid – N arbitrary 3-D points
Position index	Integer $(r, c) \in \mathbb{N}^2$	Continuous $(x, y, z) \in \mathbb{R}^3$
Pos encoding	<code>nn.Embedding(H*W, d)</code> (lookup table)	<code>MLP: $\mathbb{R}^3 \rightarrow \mathbb{R}^d$</code>
Coord range	$r \in [0, H)$, $c \in [0, W)$	<code>xyz</code> normalized to unit sphere
Generalizes to new sizes?	No (table has fixed entries)	Yes (continuous function)

A lookup table can't work for continuous coordinates – you'd need infinitely many entries. Sinusoidal features would work, but the authors chose a small MLP: **128 hidden** → **GELU** → **768**, ~100k params, learned end-to-end.

The MLP is expressive enough to **learn its own positional basis** – whatever frequency spectrum best suits the training distribution, rather than fixing it via sinusoidal priors.

```
# Learnable CLS and its own positional embedding (prepended to all sequences)
self.cls_token = nn.Parameter(torch.zeros(1, 1, 768))
self.cls_pos = nn.Parameter(torch.randn(1, 1, 768))

def forward(self, pts, colors):
    _, center, features = self.group_divider(pts, colors) # Stage 1
    tokens = self.encoder(features) # Stage 2, (B, 512, 256)
    tokens = self.encoder2trans(tokens) # Linear 256→768

    # Stage 3: CLS + pos
    cls_tokens = self.cls_token.expand(B, -1, -1)
    cls_pos = self.cls_pos.expand(B, -1, -1)
    pos = self.pos_embed(center) # (B, 512, 768)
    x = torch.cat([cls_tokens, tokens], dim=1) # (B, 513, 768)
    pos = torch.cat([cls_pos, pos], dim=1) # (B, 513, 768)
    x = x + pos # same as standard ViT

    # Stage 4: stock ViT blocks (loaded unchanged from EVA / CLIP-ViT)
    x = self.patch_dropout(x) # optional CLIP-LAION trick
    for blk in self.visual.blocks:
        x = blk(x)
    x = self.visual.norm(x[:, 0, :]) # CLS token out

    return self.trans2embed(x) # Linear 768 → 512 (CLIP dim)
```

Trainable vs Reused

What's **new for 3-D**:

- `Group` module (FPS + kNN) – **no parameters**
- `Encoder` (residual PointNet, `~1.1 M` params for $6 \rightarrow 128 \rightarrow \dots \rightarrow 256$)
- `pos_embed` MLP ($3 \rightarrow 128 \rightarrow 768$, `~100 k` params)
- `encoder2trans` (`Linear 256  $\rightarrow$  768`)
- `cls_token` , `cls_pos` (two learned vectors)
- `trans2embed` (`Linear 768  $\rightarrow$  512`)

What's **reused from 2-D ViT** (the bulk of the params):

- All transformer blocks (`visual.blocks`) – e.g. 24 blocks $\times$ `~7 M` params for a ViT-L
- Final `norm` + `fc_norm`

This is the whole point. At `~1 B` params, over 99 % of the weights come from the pretrained 2-D ViT – the 3-D adapter is tiny. Training on 10 M images + 70 M texts + 1 M shapes is affordable only because you're *fine-tuning*, not training from scratch.

Key Numbers (defaults)

Knob	Default	Role
N (input points per cloud)	~10 000	raw density
num_group	512	number of patch tokens
group_size	32	kNN size per patch
encoder_dim	256	patchifier output width
trans_dim	768 (or 1024 / 1408)	ViT width (matches backbone)
embed_dim	512 (or 768 / 1024)	CLIP alignment dim
pos_embed hidden	128	MLP hidden layer
patch_dropout	0 or 0.5	CLIP-LAION regularization

All knobs live in the `args` dict passed to `PointcloudEncoder.__init__`. The token count `num_group = 512` matters most: it sets the ViT sequence length (plus CLS → 513), which governs compute and memory.

Mental Model in One Sentence

Uni3D says:

*"Patchify a point cloud into **512 local neighborhoods**, summarize each one into a **translation-invariant token** with a mini-PointNet, and let a **continuous MLP** turn each centroid's 3-D coordinate into a **positional embedding**."*

*Then run the **same ViT weights** you'd use on an image, and align the **CLS** token to CLIP."*

Everything in that sentence is the design. The heavy lifting (transformer blocks, optimizer dynamics, alignment losses) is standard ViT + CLIP. The 3-D engineering is three modules – **Group, Encoder, pos_embed** – totaling maybe 1 % of the parameters but doing all of the geometric reasoning.

Appendix C

Xiaomi-Robotics-0

arXiv 2602.12684 · 4.7B-parameter VLA · open-sourced
A walkthrough through the lens of our Unified VLA design.

Paper at a Glance

- **Problem.** VLA models keep getting bigger and slower. At billions of params, **inference latency on real robots becomes the bottleneck** – synchronous rollouts pause the robot until the next chunk arrives, jerky transitions push the policy out of distribution.
- **Key idea.** Build a 4.7B VLA (frozen Qwen3-VL-4B + 16-layer DiT) and engineer the *post-training + deployment* pipeline for **smooth asynchronous execution**, not just headline benchmark numbers.
- **Headline numbers.** SOTA on LIBERO **98.7 %** avg, CALVIN **4.80 / 4.75** avg-len, SimplerEnv Google-Robot **85.5 % / 74.7 %**, WidowX **79.2 %**. Inference latency **80 ms on a 4090**. On real bimanual Lego/Towel tasks beats $\pi_{0.5}$ throughput.
- **What's actually new.** Three post-training tricks: action-prefix conditioning, **Λ -shape attention mask**, RoPE offset on noisy action tokens, plus dynamic loss reweighting. Plus careful 30 Hz timeline alignment at deployment.

Architecture

VLM: Qwen3-VL-4B-Instruct (text 36 layers, hidden 2560, GQA 8 KV)
Inputs: K camera images + language instruction
Outputs: KV cache (used by DiT only) + (optional) text

↓ KV cache (last 16 layers only)

DiT: 16 layers, hidden 1024, head_dim 128, 8 KV heads (GQA)
Tokens: [SINK, state_token, noisy_action_t1 ... noisy_T]
Per layer: adaLN-zero (timestep $\rightarrow$ 6 mod params)
 attention(self + cross to VLM K/V)
 ActionMLP (gated SiLU)
Causal self-attention; full attention to VLM cache

↓ output of last action_T positions

flow-matching velocity $\rightarrow$ action chunk

Mixture-of-Transformers: VLM and DiT are separate stacks; DiT cross-attends to the VLM's per-layer K/V cache (last 16 only) – no shared self-attention. $T = 30$ (1 s of actions on real robots).

Pre-training Dataset (~280M samples)

Robot trajectories (~200M timesteps)

- DROID, MolmoAct (open-sourced, multi-embodiment teleop)
- In-house **Lego Disassembly** – 338 hours of bimanual teleop
- In-house **Towel Folding** – 400 hours of bimanual teleop
- Cross-embodiment, single + dual arm; both wrist cameras and external scene cams

Vision-Language data (~80M samples)

- General VL: standard captioning / QA / grounding corpora
- VL derived **from robot trajectories** (re-labeled) so the VLM keeps embodied perception
- Four task families:
 - Visual grounding – Grounded SAM + Grounding DINO 1.5 + LLMdet **cross-validated consensus** for pixel-precise boxes
 - VQA + captioning – re-labeled by SOTA VLMs
 - Embodied reasoning & planning – EQA, high-level subtask plans, point-trajectory prediction (auto-generated from trajectory data)

Sample mixing: VL : robot = **1 : 6** during co-training (Stage 1).

Post-training Dataset

Per-robot specialization – post-training drops the cross-embodiment mixture and trains only on the target robot’s trajectories.

Real-robot task	Hours	Steps	Async batch size
Lego Disassembly (large-assembly + multi-assembly)	338 h	40k	2,048
Towel Folding	400 h	80k	2,048

Bimanual hardware: two 6-DoF arms, **3 cameras** (two wrist + one external), 30 Hz timeline.

Sim-only post-training uses the matching benchmark dataset only:

- LIBERO four splits ($T = 10$) → fine-tune from pre-trained checkpoint
- CALVIN ABCD→D and ABC→D ($T = 10$)
- SimplerEnv: RT-1 Fractal (Google Robot, $T = 4$) and Bridge (WidowX, $T = 4$)

Optimizer: **AdamW + DeepSpeed ZeRO-2** in both stages.

Input / Prompt Format

The released inference call (from `Xiaomi-Robotics-0/README.md`) shows the exact format:

```
instruction = (  
    "<im_start>user\\n"  
    "The following observations are captured from multiple views.\\n"  
    "# Base View\\n<vision_start> <image_pad> <vision_end>\\n"  
    "# Left-Wrist View\\n<vision_start> <image_pad> <vision_end>\\n"  
    "Generate robot actions for the task:\\n"  
    f"{language_instruction} /no_cot<im_end>\\n"  
    "<im_start>assistant\\n<cot></cot><im_end>\\n"  
)
```

- **No proprio in the text** – `state` is fed directly to the DiT via a 2-layer MLP (`MLPProjector`).
- **Per-camera labels** (`# Base View` , `# Left-Wrist View`) are plain text – no learned wrist delimiter.
- `/no_cot` flag toggles CoT generation; assistant turn already contains the empty `<cot></cot>` placeholder.
- **Robot-type identity** flows through `processor.get_action_mask(robot_type)` (per-dim std-mask) and `decode_action(robot_type)` (mean/std denorm) – *not* a learned embodiment token.

Pre-training Stage 1 – Choice Policies

Goal: equip the VLM itself with action-prediction capability without touching the DiT yet.

Architecture for this stage (different from final inference): append T learnable action tokens $[A_i]$ and one score token $[S]$ to the VLM input:

```
o_t, l, s_t, [A_1], [A_2], ..., [A_T], [S]
```

Each $[A_i]$ head outputs N **candidate** actions for timestep i . $[S]$ outputs N candidate scores.

Winner-takes-all (Choice Policies, Tom et al.):

1. Compute L1 distance between each candidate chunk and the GT chunk.
2. **Only the candidate with the lowest L1** receives backprop on action prediction.
3. The score head is supervised with L1 distances as targets (regression).

Why: trajectories are multimodal – many valid ways to fold a towel. A single-mode regressor averages them into mush. N candidates let the VLM represent multiple modes; the score head learns which to pick.

Pre-training Stage 2 – Flow Matching on the DiT

Freeze the VLM (its action heads from Stage 1 are dropped – VLM goes back to seeing only `o_t, l`). Train the DiT from scratch on the full robot trajectory mixture.

The flow-matching loss (Eq. 1):

$$\mathcal{L}(\theta) = \|v_{\theta}(o_t, l, s_t, \tilde{a}_{t:t+T}^{\tau}, \tau) - u(\tilde{a}_{t:t+T}^{\tau}, a_{t:t+T}, \tau)\|_2^2$$

with $\tilde{a}_{t:t+T}^{\tau} = \tau \cdot a_{t:t+T} + (1 - \tau) \epsilon, \epsilon \sim \mathcal{N}(0, I), \tau \in [0, 0.999]$.

- **Velocity target** $u = a - \epsilon$ (same as everyone else's rectified flow).
- **Beta distribution for τ sampling** – places more weight on noisier timesteps (so the network learns to "uncrap" hard cases first), following $\pi_0 / \pi_{0.5}$.
- **adaLN-zero** in every DiT layer carries the timestep $\rightarrow$ `(shift_msa, scale_msa, gate_msa, shift_mlp, scale_mlp, gate_mlp)` – six modulation knobs per block.
- **SINK + state + action** in the DiT: a learnable attention sink stabilizes attention distributions; state is one MLP-projected token; action chunk is T noised tokens.

Post-training (Sync) – Just More of the Same

- Drop the cross-embodiment mixture, train only on the target robot's trajectories.
- Unfreeze both VLM and DiT – full-model fine-tune.
- Loss is identical to Pre-training Stage 2 (flow-matching MSE).
- Deploy synchronously: roll out `T_e` steps, **stop, infer, stop, repeat**. Robot pauses each cycle.

This is the strong baseline you'd expect – full SFT on real-robot data. The interesting bit is the asynchronous variant on the next slide.

Post-training (Async) – The Real Contribution

Setup. Condition action generation on Δt_c previously committed actions, prefixed before the noisy chunk:

```
[SINK, s_t, a_t, ..., a_{t+\Delta t_c-1}, \hat{a}_{t+\Delta t_c}, ..., \hat{a}_{t+T-1}]  
      | clean prefix (committed) | | noisy actions (to predict) |
```

Three tricks to prevent "shortcut by copy":


1. **RoPE offset of +10** on noisy-action positional indices – gives the model a discrete cue that distinguishes "this came from the previous chunk" vs "I'm predicting this now".
2. **Λ -shape attention mask** (next slide) – early noisy tokens attend to the prefix for smooth stitching, but later noisy tokens **cannot see the prefix at all**, forcing them to attend to vision/language K/V.
3. **Dynamic loss reweighting**: when $\Delta t_c > 0$, weight each sample's flow-matching loss by the L1 between online-predicted actions and GT – focuses gradient on samples where the policy is *actually wrong*, not the easy ones.

Λ-Shape Attention Mask

Query \ Key:	VLM_KV	SINK	state	prefix_a_0.. $\Delta t_c - 1$	noisy_t+ Δt_c .. $T - 1$
VLM_KV	(frozen, fed via cache – DiT just reads it as keys)				
SINK	✓	self			
state	✓	✓	self		
prefix	✓	✓	✓	causal	
noisy[i<w]	✓	✓	✓	full to prefix	causal
noisy[i≥w]	✓	✓	✓	BLOCKED	causal ← Λ kink

- The "Λ" comes from the kink: for the first **w** noisy tokens (immediately after the prefix), attention to the prefix is **on**, allowing smooth transition. For all later noisy tokens, attention to the prefix is **off**, forcing them to attend to vision/language signals.
- Implementation note: the **released codebase** (`xiaomi_mibot/modeling_mibot.py`) is **inference only**. The mask in `MiBoTForActionGeneration.forward` (line 1832) is plain causal `torch.tril( ... )` – the Λ-mask lives in the (un-released) training code only.
- The shape comes from prior literature (sliding-window, sink-attention work cited as refs 16, 20, 71).

Asynchronous Deployment



- **Inference latency** $t_{inf} \approx 80$ ms on a single 4090 (FlashAttn 2 + bf16, 5 FM steps, $T = 30$).
- Robot **never stops**: while the model infers chunk N+1, the robot continues executing the tail of chunk N. As soon as N+1 lands, it stitches in.
- Stitch point: the action prefix in chunk N+1 covers steps $[T_e, T_e + \Delta t_c)$ of the previous chunk. Set $\Delta t_c \geq \Delta t_{inf} / dt$ so the prefix bridges the entire inference window – no gaps.
- All sensors **resampled to a unified 30 Hz timeline** at deployment; nearest-timestamp aggregation per tick.

Codebase Map

What lives where in `xiaomi_mibot/` (the released HF-format inference release):

Concept	File:line	Notes
DiT (16 <code>DecoderLayer</code> s)	<code>modeling_mibot.py:1710</code>	Loops over layers; cross-attends to last 16 of past_key_values
<code>DecoderLayer</code> (adaLN-zero)	<code>modeling_mibot.py:1676</code>	<code>ada_ln_table[None] + t_embeds</code> → 6 mod params, then attn + ActionMLP
<code>Attention</code> (self + cross via cat)	<code>modeling_mibot.py:1611</code>	Concat VLM K/V (cache) with current K/V before SDPA
Sink token	<code>modeling_mibot.py:1782</code>	<code>nn.Embedding(1, hidden_size)</code>
State token (MLP)	<code>modeling_mibot.py:1754</code>	2-layer MLP <code>state_dim=32</code> → <code>hidden=1024</code>
Action input projector	<code>modeling_mibot.py:1759</code>	2-layer MLP <code>action_dim=32</code> → <code>1024</code> , applied after <code>noisy_action *= action_mask</code>
Action output projector	<code>modeling_mibot.py:1764</code>	Same shape, mirrors input
FM sampling loop	<code>modeling_mibot.py:1868–1871</code>	<code>x = x + v * dt</code> , 5 Euler steps
Position IDs (DiT cont. from VLM)	<code>modeling_mibot.py:1824</code>	<code>arange(...) + vlm.position_ids.max() + 1</code>
Causal mask + cache mask	<code>modeling_mibot.py:1832–1834</code>	Plain <code>tril</code> ; $\wedge$ -mask is training-only and not released
Action mask (per robot)	<code>processing_mibot.py:118</code>	<code>(action_config[robot_type]["std"] > 1e-5).float()</code>
Action denorm	<code>processing_mibot.py:132</code>	<code>actions * std + mean</code>
Prompt template	<code>Xiaomi-Robotics-0/README.md:104</code>	Plain text with <code># Base View</code> / <code># Left-Wrist View</code> labels

Compare & Contrast – Architecture

Axis	Xiaomi-Robotics-0	Unified VLA (ours)
VLM	Qwen3-VL-4B-Instruct (4B, frozen in Stage 2)	Qwen2.5-VL-3B (3B, frozen Phase 1)
Action stack	Separate 16-layer DiT (1024 hidden, 8 KV) outside the VLM	In-place ExpertBlocks (PC + Action) at every Qwen layer (or sparse subset)
Cross-attn topology	DiT layer $i \leftrightarrow$ VLM layer (last 16 only) K/V cache	Each expert layer $\leftrightarrow$ same-layer Qwen K/V; Action also reads same-layer PC K/V
Modalities	Image + text (no point cloud)	Image + text + point cloud
Proprio	Single MLP-projected token	$P \times \langle \text{proprio_start} \rangle \text{ProprioEncoder} \langle \text{proprio_end} \rangle$; variable P (current or current+history)
Embodiment	Implicit (per-robot post-training) + per-dim action-mask + mean/std denorm	SharedEmbodimentEmbedding : shared bottleneck $\rightarrow$ embodiment_pc + embodiment_action
Sink token	$\checkmark$ <code>nn.Embedding(1, hidden)</code> at front of DiT	Not explicit; embodiment_action plays a similar prefix-bias role
Timestep injection	adaLN-zero (6 mod params + zero-init gates)	adaLN-zero (6 mod params + zero-init gates), timestep-only – matches DiT exactly
Parameter total	4.7B	$\sim$ 3B VLM + small experts (varies with placement)

Compare & Contrast – Training & Deployment

Axis	Xiaomi-Robotics-0	Unified VLA (ours)
Training stages	(1) VLM with action heads (Choice Policies WTA + VL NTP) → (2) freeze VLM, train DiT (FM) → (3) per-robot SFT	Single-stage Phase 1: VLM frozen, experts trainable from start
Multimodality handling	N candidates + WTA + score head	Standard FM regression to one velocity (no multi-mode)
Pretraining data	200 M robot timesteps + 80 M VL (1:6)	(Dataset not yet wired up in repo)
Catastrophic-forgetting guard	VL co-training + Stage 2 freeze	VLM stays frozen → forgetting is structurally impossible
Loss flow	Stage 2: action-loss reaches DiT only (VLM K/V cache is detached implicitly)	Action-loss reaches PC and Action experts (PC K/V not detached); VLM K/V detached
Async / RTC	Action-prefix conditioning + A-mask + RoPE offset + dynamic loss reweighting	None – synchronous only
Inference latency	80 ms / chunk @ 4090 (T=30, 5 FM steps)	Untested (no inference-with-cache loop yet)
Sensor sync	30 Hz unified timeline at deployment	Untested

Improvements Worth Stealing – Tier 1 (low cost, high value)

1. Λ -shape attention mask + RoPE offset for action prefix

Adds async support directly. Fits into our existing causal Action-segment mask: split current causal triangle into `[full_to_prefix | causal_self]` for early noisy tokens, `[blocked | causal_self]` for later. RoPE offset is a $+N$ on the integer positions of noisy action tokens before we feed them into MRoPE. **Estimated effort: small Action-mask change + dataloader for prefix.**

2. Dynamic loss reweighting by L1 error

Drop-in change to `flow_matching_loss`: weight each sample by L1 between current prediction and GT. Focuses gradient on hard cases. **Estimated effort: ~10 lines.**

3. β -distribution for τ sampling

We currently sample τ uniformly. Switching to $\text{Beta}(\alpha, \beta)$ with $\alpha < \beta$ places weight on noisier $\tau \rightarrow$ better convergence on hard timesteps. Same change as π_0 / Wan / Qwen-Image. **Estimated effort: 1 line.**

Improvements Worth Stealing – Tier 2 (architectural)

4. SINK token DONE

We now have separate `sink_pc` and `sink_action` Parameters at position 0 of each segment.

5. Two-stage pretraining (Choice Policies → frozen VLM + DiT)

The Stage-1 winner-takes-all multi-candidate scheme directly addresses **trajectory multimodality** – something our single-velocity FM regressor doesn't handle. Even if we keep our backbone-multi-expert architecture, we could add an N-candidate + score head during a Stage-1 warmup, then keep only the FM training in Stage 2.

6. adaLN-zero with explicit gates DONE

We now produce 6 mod params per ExpertBlock `(s1, sh1, g1, s2, sh2, g2)`; `modulate` uses the `(1 + scale)` form; final Linear of both adaLN MLPs is zero-initialized → block is residual identity at step 0.

Improvements Worth Stealing – Tier 3 (deployment / async)

8. 30 Hz unified-timeline sensor sync

At deployment, resample images / proprio / point clouds onto a single clock and feed nearest-timestamp values per tick. Avoids the silent latency drift when sensors run at different rates.

9. Last-N VLM layers only for K/V cache

They use only the **last 16** VLM layers' K/V instead of all 36. This is essentially the same idea as our `expert_layer_ids` (sparse expert placement) but applied from the VLM side. We could similarly cap which Qwen K/V the experts read.

10. Action-mask on input dimensions (per-robot active-DoF)

Multiply noisy actions by a `(1, 1, action_dim)` mask before projection – zeros out irrelevant DoFs for that robot and prevents gradient on inactive dims. Useful when we mix robots with different DoF (Trossen 7 / Franka 7 / SO-101 6) inside one batch.

What We Already Do Better

- **Cross-embodiment learning is explicit** – `SharedEmbodimentEmbedding` is a single shared bottleneck for all robots in PC and Action streams; their per-robot post-training discards cross-embodiment knowledge by construction. Ours can in principle handle a new robot zero-shot if it gets a new ID; theirs can't.
- **Point clouds as a first-class modality** – Xiaomi only uses RGB; our PC expert reads geometry directly via Uni3D + a dedicated bidirectional expert with cross-attention to VLM K/V. For tasks where 3-D geometry matters (insertion, stacking), this should give us a structural advantage.
- **In-place expert design** – our PC and Action experts sit *next to* every Qwen layer and are bit-for-bit verified to leave VLM forward unchanged. Theirs has a clean modular split but pays for it: their DiT can only consume the VLM's *post-hoc* representations. We get full per-layer cross-modal mixing.
- **No detach on PC → Action** – action loss flows back into the PC expert through cross-attention, so PC representations align to action quality, not just point-cloud reconstruction quality.
- **Variable proprio history** – `<proprio_start> P × ProprioEncoder <proprio_end>` handles current-only or current + history with the same Linear; theirs is one fixed state token.

Suggested Roadmap

Now (Phase 1):

1. Λ -mask + RoPE offset + dynamic loss reweighting (Tier 1) – unlocks async deployment for our pipeline too.
2. β -distribution τ sampling – single-line change for the FM loss.
3. SINK token at front of Action segment.

Phase 2 (post-VLM-unfreezing): 4. VL co-training mix (1:6) to avoid catastrophic forgetting. 5. Two-stage pretraining curriculum (Choice Policies WTA $\rightarrow$ FM-only on DiT) **if** trajectory multimodality starts hurting our policies.

Deployment: 6. 30 Hz unified-sensor timeline. 7. Action-mask on input DoFs per robot type.

The biggest near-term win is async post-training. Their headline result against $n_{0.5}$ comes from those three tricks, not from raw model size.